

The files and their contents are described below:

original_dataset.csv: Contains the original data collected from Scopus, consisting of 15,488 records and 20 columns. The column "abstract_text" contains the abstract text, the column "Sustainable Development Goals (2023)" contains the SDG as a string, columns SDG 1-17 contain the SDGs in a one-hot encoded format, and the column "label" contains the unique SDGs.

preprocessed_dataset.csv: Contains the dataset after preprocessing. The column "processed_abstract_text" contains the paper abstracts, and the column "label" contains the respective SDG.

original_training_set.csv: Contains 80% of the original unprocessed data for training, after a stratified 80/20 split. The columns include the abstract text and the respective SDG label.

original_testing_set.csv: Contains 20% of the original unprocessed data set aside for testing, after a stratified 80/20 split. The columns include the abstract text and the respective SDG label.

balanced_training_set.csv: Contains the training set after preprocessing and balancing, with 300 samples per class. Column "abstract" contains the abstract text, column "label" contains the SDG label, and column "synthetic" contains a boolean indicating whether the sample was generated by a generative model. Samples marked with True in column "synthetic", were generated.